

Refine System Requirements

Minimum hardware and configuration requirements for running Refine at different scales.

Quick Reference

	Local (1-5 accounts)	Starter (1-10 accounts)	Business (10-50 accounts)	Enterprise (50-200 accounts)	Large Enterprise (200+)
Where	Your laptop/desktop	EC2 <code>t3.small</code>	EC2 <code>t3.medium</code>	EC2 <code>m6i.xlarge</code>	EC2 <code>m6i.2xlarge</code>
CPU	Any modern CPU	1 vCPU	2 vCPUs	4 vCPUs	8+ vCPUs
RAM	2 GB (Docker Desktop)	512 MB	1-2 GB	2-4 GB	4-8 GB
Storage	1 GB	1 GB	5 GB	20 GB	50+ GB
Database	SQLite (included)	SQLite (included)	SQLite (WAL mode)	PostgreSQL recommended	PostgreSQL required
Dashboard Load	< 1 second	< 1 second	1-2 seconds	2-4 seconds	Requires pagination
Nightly Sync	< 3 minutes	< 5 minutes	5-15 minutes	15-45 minutes	Requires parallelization
Always-on	No (sleeps with laptop)	Yes	Yes	Yes	Yes

Multi-User Overhead

Refine uses a shared server model — all users connect to the same Refine instance. Each additional concurrent user adds minimal overhead (approximately 10-20 MB of RAM per active session). The resource requirements below are sufficient for typical multi-user workloads. No additional scaling is needed for user count alone; the primary scaling factor remains the number of monitored accounts — AWS accounts, Azure subscriptions, and GCP projects each count toward the scale tiers below.

Running Locally (1-5 AWS Accounts)

For: Evaluation, demos, small teams, personal use.

Refine runs on any Windows, Mac, or Linux machine with Docker Desktop installed. This is the fastest way to get started — extract the ZIP, run the setup script, and you're live in under 5 minutes.

Requirements: Docker Desktop (2 GB RAM allocated) + Python 3.8+

Limitations:

- Nightly sync only runs when your machine is on — data goes stale if the laptop sleeps overnight
- Other applications compete for CPU/RAM — close heavy apps during use
- Only accessible at `localhost:8000` (teammates can't access it unless you share your local network)
- Performance degrades above 5 AWS accounts as data volume grows

When to move to EC2: When you need always-on monitoring, team access, or connect more than 5 AWS accounts.

See the [Deployment Guide](#) for setup instructions.

Tier 1: Starter (1-10 AWS Accounts)

Typical customer: Small startup or single-team environment.

Minimum Requirements

Resource	Minimum	Recommended
CPU	1 vCPU	2 vCPUs
RAM	512 MB	1 GB
Disk	1 GB	2 GB
OS	Any Linux (Docker-capable) or Windows with WSL2	Ubuntu 22.04 LTS
Docker	Docker Engine 20.10+	Docker Engine 24+
Network	Outbound HTTPS to AWS, Azure & GCP APIs	Same

Database

SQLite (default, included) is sufficient. No additional setup required.

Expected Performance

- **Dashboard load:** < 500ms
- **API response time:** < 200ms average
- **Nightly sync duration:** 2-5 minutes
- **Data volume:** ~600 KB/account/month of inventory data
- **Database size:** 20-30 MB after 12 months

Suitable Cloud Instances

Provider	Instance Type	Monthly Cost
AWS	t3.small (2 vCPU, 2 GB)	~\$15/mo
Azure	B1ms (1 vCPU, 2 GB)	~\$15/mo
GCP	e2-small (2 vCPU, 2 GB)	~\$14/mo
On-premises	Any machine with Docker	\$0 (hardware cost only)

Tier 2: Business (10-50 AWS Accounts)

Typical customer: Mid-size company with separate accounts for production, staging, and development per team.

Minimum Requirements

Resource	Minimum	Recommended
CPU	2 vCPUs	4 vCPUs
RAM	1 GB	2 GB
Disk	5 GB	10 GB
OS	Any Linux (Docker-capable)	Ubuntu 22.04 LTS
Docker	Docker Engine 20.10+	Docker Engine 24+
Network	Outbound HTTPS to AWS, Azure & GCP APIs	Same

Database

SQLite is still workable at this tier but may show occasional slowness during nightly sync when many accounts are updating simultaneously.

Recommended: Enable SQLite WAL (Write-Ahead Logging) mode for better concurrent read/write performance. Set this in your `.env`:

```
DATABASE_URL=sqlite:///app/data/refine.db?mode=wal
```

Optional upgrade: PostgreSQL provides better concurrent access and query performance at this scale. If you choose PostgreSQL:

```
DATABASE_URL=postgresql://refine:password@localhost:5432/refine
```

Expected Performance

- **Dashboard load:** 1-2 seconds

- **API response time:** 200-500ms average
- **Nightly sync duration:** 5-15 minutes
- **Data volume:** ~30 MB/month total across all accounts
- **Database size:** 75-150 MB after 12 months
- **Container memory usage:** 400 MB - 1 GB during peak API loads

Suitable Cloud Instances

Provider	Instance Type	Monthly Cost
AWS	t3.medium (2 vCPU, 4 GB)	~\$30/mo
Azure	B2s (2 vCPU, 4 GB)	~\$30/mo
GCP	e2-medium (2 vCPU, 4 GB)	~\$27/mo

Tier 3: Enterprise (50-200 AWS Accounts)

Typical customer: Large organization with multi-account strategy (AWS Organizations), multiple business units, and shared services accounts.

Minimum Requirements

Resource	Minimum	Recommended
CPU	4 vCPUs	8 vCPUs
RAM	2 GB	4 GB
Disk	20 GB SSD	50 GB SSD
OS	Linux (Ubuntu 22.04 LTS recommended)	Same
Docker	Docker Engine 24+	Same
Database	PostgreSQL 14+ (separate)	PostgreSQL 16 with connection pooling
Network	Outbound HTTPS (AWS, Azure & GCP) + internal DB access	Same

Database

PostgreSQL is strongly recommended at this tier. SQLite's single-writer limitation causes sync bottlenecks and API timeouts when more than 50 accounts are actively writing data.

```
# docker-compose.yml addition for PostgreSQL
services:
  postgres:
    image: postgres:16
    environment:
      POSTGRES_DB: refine
      POSTGRES_USER: refine
      POSTGRES_PASSWORD: your-secure-password
    volumes:
      - postgres_data:/var/lib/postgresql/data
    deploy:
      resources:
        limits:
          memory: 1G

  refine:
    image: public.ecr.aws/a4k0g9y9/refine:latest
    environment:
      DATABASE_URL: postgresql://refine:your-secure-password@postgres:5432/refine
    depends_on:
      - postgres
    deploy:
      resources:
        limits:
          memory: 4G
          cpus: "4"
```

Expected Performance

- **Dashboard load:** 2-4 seconds (large datasets)
- **API response time:** 500ms - 2 seconds for inventory endpoints
- **Nightly sync duration:** 15-45 minutes (sequential account processing)
- **Data volume:** ~100-120 MB/month total
- **Database size:** 200-400 MB after 12 months
- **Container memory usage:** 1-3 GB during peak loads (dashboard + savings recalculation)

Performance Considerations

At this tier, you may notice:

1. **Dashboard load times increase** — The dashboard loads all 14 service inventories in parallel. With 200 accounts and ~60 EC2 instances each, that's 12,000+ instances loaded into memory per API call.
2. **Nightly sync takes longer** — Each account is synced sequentially. With 200 accounts averaging 10 minutes each, the full sync cycle takes 30-45 minutes.
3. **Savings recalculation is heavier** — The savings engine iterates through every EC2 and RDS instance to calculate rightsizing recommendations. At 12,000+ instances, this takes 1-2 seconds per call.

Suitable Cloud Instances

Provider	Instance Type	Monthly Cost
AWS	m6i.xlarge (4 vCPU, 16 GB)	~\$140/mo
Azure	D4s v5 (4 vCPU, 16 GB)	~\$140/mo
GCP	e2-standard-4 (4 vCPU, 16 GB)	~\$120/mo

Note: Memory is the primary scaling constraint, not CPU. Choose instances with more RAM over more cores.

Tier 4: Large Enterprise (200+ AWS Accounts)

Typical customer: Large enterprise with AWS Organizations, hundreds of accounts across multiple regions and business units.

Minimum Requirements

Resource	Minimum	Recommended
CPU	8 vCPUs	16 vCPUs (across replicas)
RAM	4 GB	8 GB per replica
Disk	50 GB SSD	100 GB SSD
OS	Linux (Ubuntu 22.04 LTS)	Same
Docker	Docker Engine 24+ or Kubernetes	Kubernetes recommended
Database	PostgreSQL 16 with connection pooling	PostgreSQL 16 + pgBouncer
Cache	Redis 7+ (optional but recommended)	Redis 7+
Network	Outbound HTTPS + internal DB/cache access	Same

Architecture Recommendations

At 200+ accounts, a single Refine container is no longer sufficient. We recommend:

1. **Multiple Refine replicas** behind a load balancer (2-3 containers)
2. **PostgreSQL** as the primary database with connection pooling
3. **Redis** for rate limiting and API response caching
4. **Separate sync worker** — run the nightly sync in a dedicated container so it doesn't compete with API traffic

```
# docker-compose.yml for large-scale deployment
services:
  postgres:
    image: postgres:16
    environment:
      POSTGRES_DB: refine
      POSTGRES_USER: refine
      POSTGRES_PASSWORD: ${DB_PASSWORD}
    volumes:
      - postgres_data:/var/lib/postgresql/data
    deploy:
      resources:
        limits:
          memory: 2G

  redis:
    image: redis:7-alpine
    deploy:
      resources:
        limits:
          memory: 256M

  refine-api:
    image: public.ecr.aws/a4k0g9y9/refine:latest
    environment:
      DATABASE_URL: postgresql://refine:${DB_PASSWORD}@postgres:5432/refine
      RATE_LIMIT_STORAGE_URI: redis://redis:6379
      DISABLE_SCHEDULER: "true" # API replicas don't run sync
    deploy:
      replicas: 2
      resources:
        limits:
          memory: 4G
          cpus: "4"

  refine-sync:
    image: public.ecr.aws/a4k0g9y9/refine:latest
    environment:
      DATABASE_URL: postgresql://refine:${DB_PASSWORD}@postgres:5432/refine
      SYNC_ONLY: "true" # Only runs scheduler, no API
    deploy:
      replicas: 1
      resources:
        limits:
          memory: 4G
          cpus: "2"
```

Expected Performance

- **Dashboard load:** 3-6 seconds (consider using account filtering)
- **API response time:** 1-3 seconds for large inventory queries
- **Nightly sync duration:** 45+ minutes (sequential); can be reduced with parallel sync
- **Data volume:** 200+ MB/month
- **Database size:** 500 MB - 1 GB after 12 months

Scaling Strategies

1. **Use account filtering** — Instead of loading all 200+ accounts on the dashboard, select specific accounts or account groups. This dramatically reduces payload size and response times.
2. **Stagger nightly syncs** — Configure different sync times for different account groups to spread the load.
3. **Monitor container memory** — Set Docker memory limits and monitor usage. Memory spikes during dashboard loads and savings recalculations are the primary concern.

Suitable Cloud Instances

Provider	Instance Type	Monthly Cost (2 replicas)
AWS	m6i.2xlarge (8 vCPU, 32 GB) x2	~\$560/mo
Azure	D8s v5 (8 vCPU, 32 GB) x2	~\$560/mo
GCP	e2-standard-8 (8 vCPU, 32 GB) x2	~\$480/mo
AWS ECS/Fargate	4 vCPU, 8 GB x2	~\$400/mo

Add ~\$50-100/mo for managed PostgreSQL (RDS) and ~\$15/mo for ElastiCache Redis.

Storage Sizing Guide

Disk Space by Scale

Component	10 Accounts	50 Accounts	200 Accounts
Docker image	500 MB	500 MB	500 MB
SQLite/PostgreSQL	30 MB	150 MB	400 MB
Inventory JSON (12 months)	70 MB	350 MB	1.4 GB
Cost summaries (12 months)	2 MB	10 MB	40 MB
Savings report data	5 MB	25 MB	100 MB
Application logs	50 MB	200 MB	500 MB
Total	~660 MB	~1.2 GB	~3 GB

Recommendation: Provision at least 3x the expected total to allow for growth and temporary files.

Network Requirements

Outbound Access

Refine needs outbound HTTPS (port 443) access to:

Endpoint	Purpose
<code>sts.amazonaws.com</code>	Cross-account role assumption
<code>lambda.{region}.amazonaws.com</code>	Invoke data collection Lambda
<code>s3.{region}.amazonaws.com</code>	Read inventory data from customer S3 buckets
<code>public.ecr.aws</code>	Pull Docker image updates (commercial deployments)

Inbound Access

Port	Purpose
8000 (configurable)	Refine web interface

For air-gapped/government deployments, see [AIR GAPPED SETUP GUIDE.md](#).

Monitoring Recommendations

What to Monitor

Metric	Warning Threshold	Critical Threshold
Container memory usage	> 70% of limit	> 90% of limit
API response time (p95)	> 2 seconds	> 5 seconds
Nightly sync duration	> 30 minutes	> 60 minutes
Database size	> 500 MB (SQLite)	> 1 GB (SQLite)
Disk usage	> 70%	> 90%

Health Check

Refine exposes a health endpoint at `GET /health` that returns `{"status": "ok"}`. Use this for Docker health checks, load balancer probes, or monitoring tools.

FAQ

Q: Can I run Refine on a Raspberry Pi? A: Only for 1-3 accounts. The ARM-based Pi 4 with 4 GB RAM can run the Docker image, but performance will be limited.

Q: Do I need a GPU? A: No. Refine is CPU and memory bound, not GPU bound.

Q: Can I use an existing PostgreSQL instance? A: Yes. Set `DATABASE_URL` in your `.env` file to point to any PostgreSQL 14+ server. Refine creates its own tables on first startup.

Q: What happens if the container runs out of memory? A: Docker will kill the container (OOMKilled). It will auto-restart if `restart: unless-stopped` is set. Consider increasing the memory limit or upgrading to the next tier.

Q: Can I run multiple Refine instances against the same database? A: Yes, with PostgreSQL. Do not use SQLite with multiple instances. Ensure only one instance runs the nightly sync scheduler (set `DISABLE_SCHEDULER=true` on API-only replicas).

Q: How do I migrate from SQLite to PostgreSQL? A: Export your data with `sqlite3 data/refine.db .dump > backup.sql`, create a PostgreSQL database, update `DATABASE_URL` in `.env`, and restart Refine. Tables are auto-created. Historical data requires manual migration.